

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import math
import cv2
from glob import glob
import itertools

from keras.utils import np_utils
from keras.preprocessing.image import ImageDataGenerator
from keras.models import Sequential
from keras.layers import Dense
from keras.layers import Dropout
from keras.layers import Flatten
from keras.layers.convolutional import Conv2D
from keras.layers.convolutional import MaxPooling2D
from keras.layers import BatchNormalization
from keras.callbacks import ModelCheckpoint, ReduceLROnPlateau, CSVLogger

from sklearn import preprocessing
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix

import tensorflow.compat.v1 as tf
tf.disable_v2_behavior()

↳ WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow_core/python
Instructions for updating:
non-resource variables are not supported in the long term

# GLOBAL VARIABLES
scale = 70
seed = 7

from google.colab import drive
drive.mount('/drive')

↳ Go to this URL in a browser: https://accounts.google.com/o/oauth2/auth?client\_id=9473

Enter your authorization code:
.....
Mounted at /drive

path_to_images = '/drive/My Drive/7jan/train/*/*.png'

# 1. Read the images and generate the train and test dataset (5 points)

images = glob(path_to_images)
trainingset = []


```

```
traininglabels = []
num = len(images)
count = 1

#READING IMAGES AND RESIZING THEM
for i in images:
    print(str(count)+'/'+str(num),end='\r')
    trainingset.append(cv2.resize(cv2.imread(i),(scale,scale)))
    traininglabels.append(i.split('/')[-2])
    count=count+1
trainingset = np.asarray(trainingset)
traininglabels = pd.DataFrame(traininglabels)
```

⇨

```
#Cleaning the images and removing the background: HSV Hue saturation value
```

Convert the RGB image into HSV.

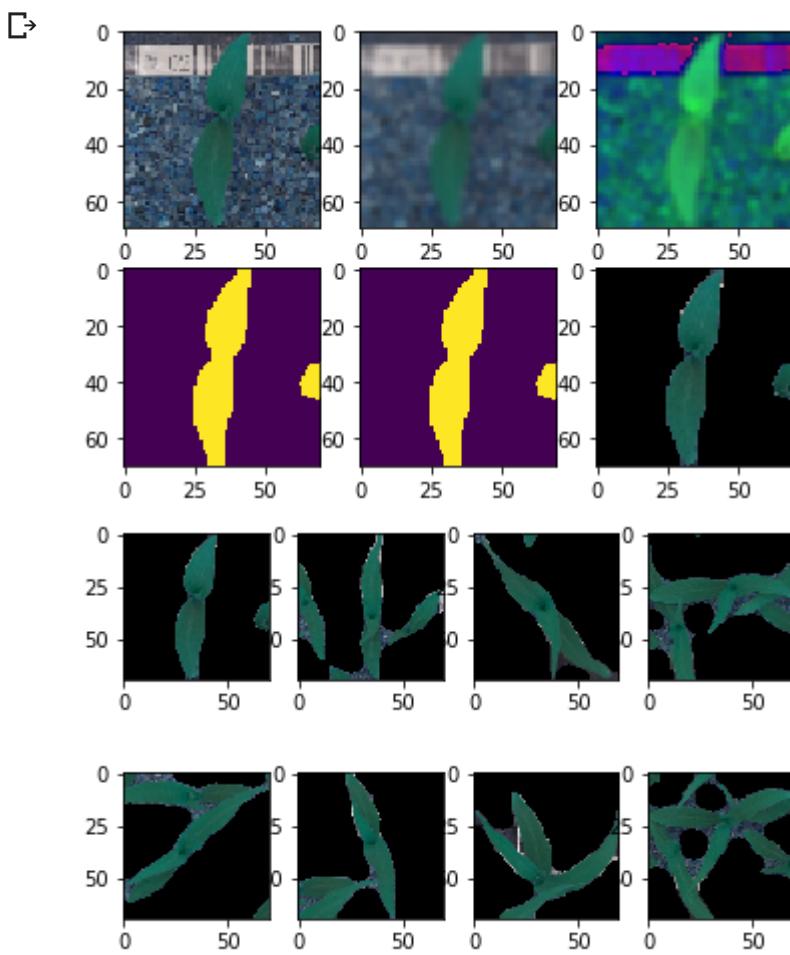
We will have to blur the image to remove noise.

We will have to create a mask to remove the background.

```
new_train = []
sets = []; getEx = True
for i in trainingset:
    blurr = cv2.GaussianBlur(i,(5,5),0)
    hsv = cv2.cvtColor(blurr,cv2.COLOR_BGR2HSV)
    #GREEN PARAMETERS
    lower = (25,40,50)
    upper = (75,255,255)
    mask = cv2.inRange(hsv,lower,upper)
    struc = cv2.getStructuringElement(cv2.MORPH_ELLIPSE,(11,11))
    mask = cv2.morphologyEx(mask,cv2.MORPH_CLOSE,struc)
    boolean = mask>0
    new = np.zeros_like(i,np.uint8)
    new[boolean] = i[boolean]
    new_train.append(new)

    if getEx:
        plt.subplot(2,3,1);plt.imshow(i) # ORIGINAL
        plt.subplot(2,3,2);plt.imshow(blurr) # BLURRED
        plt.subplot(2,3,3);plt.imshow(hsv) # HSV CONVERTED
        plt.subplot(2,3,4);plt.imshow(mask) # MASKED
        plt.subplot(2,3,5);plt.imshow(boolean) # BOOLEAN MASKED
        plt.subplot(2,3,6);plt.imshow(new) # NEW PROCESSED IMAGE
        plt.show()
    getEx = False
new_train = np.asarray(new_train)

# CLEANED IMAGES
for i in range(8):
    plt.subplot(2,4,i+1)
    plt.imshow(new_train[i])
```



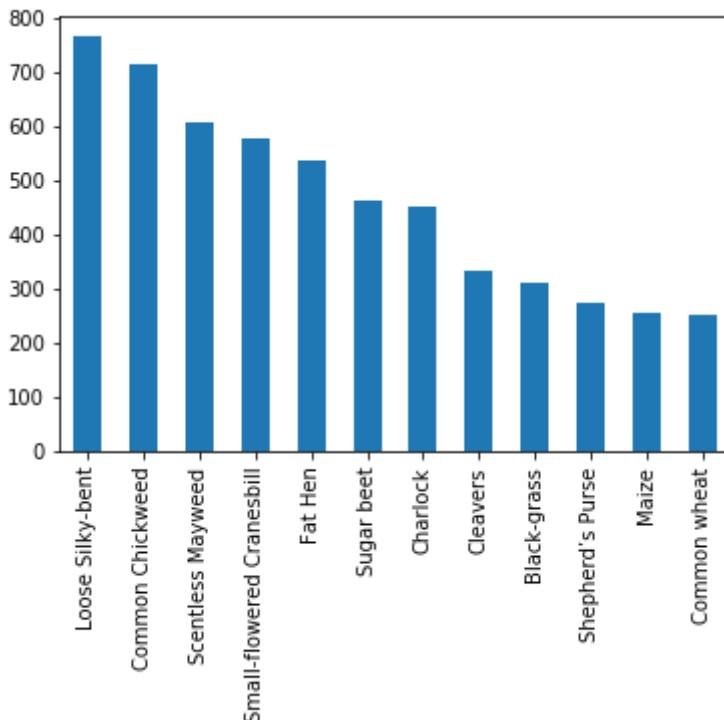
```
# convert class vectors to binary class matrices
labels = preprocessing.LabelEncoder()
labels.fit(traininglabels[0])
print('Classes'+str(labels.classes_))
encodedlabels = labels.transform(traininglabels[0])
clearalllabels = np_utils.to_categorical(encodedlabels)
classes = clearalllabels.shape[1]
print(str(classes))
traininglabels[0].value_counts().plot(kind='bar')
```

↳

```
Classes['Black-grass' 'Charlock' 'Cleavers' 'Common Chickweed' 'Common wheat'
 'Fat Hen' 'Loose Silky-bent' 'Maize' 'Scentless Mayweed'
 'Shepherd's Purse' 'Small-flowered Cranesbill' 'Sugar beet']
```

```
12
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7fa273436c88>
```



```
new_train = new_train/255
x_train,x_test,y_train,y_test = train_test_split(new_train,clearalllabels,test_size=0.1,ra
generator = ImageDataGenerator(rotation_range = 180,zoom_range = 0.1,width_shift_range = 0
generator.fit(x_train)
```

### \*3. Initialize & build the model (10 points) CNN \*

```
np.random.seed(seed)

model = Sequential()

model.add(Conv2D(filters=64, kernel_size=(5, 5), input_shape=(scale, scale, 3), activation='relu'))
model.add(BatchNormalization(axis=3))
model.add(Conv2D(filters=64, kernel_size=(5, 5), activation='relu'))
model.add(MaxPooling2D((2, 2)))
model.add(BatchNormalization(axis=3))
model.add(Dropout(0.1))

model.add(Conv2D(filters=128, kernel_size=(5, 5), activation='relu'))
model.add(BatchNormalization(axis=3))
model.add(Conv2D(filters=128, kernel_size=(5, 5), activation='relu'))
model.add(MaxPooling2D((2, 2)))
model.add(BatchNormalization(axis=3))
model.add(Dropout(0.1))

model.add(Conv2D(filters=256, kernel_size=(5, 5), activation='relu'))
model.add(BatchNormalization(axis=3))
model.add(Conv2D(filters=256, kernel_size=(5, 5), activation='relu'))
```

```
model.add(Conv2D(filters=256, kernel_size=(3, 3), activation='relu'))
model.add(BatchNormalization(axis=3))
model.add(Conv2D(filters=256, kernel_size=(5, 5), activation='relu'))
model.add(MaxPooling2D((2, 2)))
model.add(BatchNormalization(axis=3))
model.add(Dropout(0.1))

model.add(Flatten())

model.add(Dense(256, activation='relu'))
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Dense(256, activation='relu'))
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Dense(classes, activation='softmax'))

model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])

model.summary()
```



```
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorf1
Instructions for updating:
Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/optimizers.py:79
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorf1
Model: "sequential_1"
```

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 66, 66, 64)	4864
batch_normalization_1 (Batch Normalization)	(None, 66, 66, 64)	256
conv2d_2 (Conv2D)	(None, 62, 62, 64)	102464
max_pooling2d_1 (MaxPooling2D)	(None, 31, 31, 64)	0
batch_normalization_2 (Batch Normalization)	(None, 31, 31, 64)	256
dropout_1 (Dropout)	(None, 31, 31, 64)	0
conv2d_3 (Conv2D)	(None, 27, 27, 128)	204928
batch_normalization_3 (Batch Normalization)	(None, 27, 27, 128)	512
conv2d_4 (Conv2D)	(None, 23, 23, 128)	409728
max_pooling2d_2 (MaxPooling2D)	(None, 11, 11, 128)	0
batch_normalization_4 (Batch Normalization)	(None, 11, 11, 128)	512
dropout_2 (Dropout)	(None, 11, 11, 128)	0
conv2d_5 (Conv2D)	(None, 7, 7, 256)	819456

batch_normalization_5 (Batch	(None, 7, 7, 256)	1024
conv2d_6 (Conv2D)	(None, 3, 3, 256)	1638656
max_pooling2d_3 (MaxPooling2	(None, 1, 1, 256)	0
batch_normalization_6 (Batch	(None, 1, 1, 256)	1024
dropout_3 (Dropout)	(None, 1, 1, 256)	0
flatten_1 (Flatten)	(None, 256)	0
dense_1 (Dense)	(None, 256)	65792
batch_normalization_7 (Batch	(None, 256)	1024
dropout_4 (Dropout)	(None, 256)	0
dense_2 (Dense)	(None, 256)	65792
batch_normalization_8 (Batch	(None, 256)	1024
dropout_5 (Dropout)	(None, 256)	0
dense_3 (Dense)	(None, 12)	3084
<hr/>		
Total params:	3,320,396	
Trainable params:	3,317,580	
Non-trainable params:	2,816	

```

from keras.utils import np_utils # convert to one-hot-encoding
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten, Conv2D, MaxPool2D, GlobalMaxPooling2D
from keras.optimizers import RMSprop
from keras.preprocessing.image import ImageDataGenerator
from keras.optimizers import Adam

#4. Optimize the model (8 points)

opt = Adam(lr=0.001)

optimizer = RMSprop(lr=0.001, rho=0.9, epsilon=1e-08, decay=0.0)
model.compile(optimizer = optimizer, loss = "categorical_crossentropy", metrics = ["accuracy"])

print('x_train shape:', x_train.shape)
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

```

```
↳ x_train shape: (4989, 70, 70, 3)
4989 train samples
555 test samples
```

```
batch_size = 128
num_classes = 10
epochs = 12
```

```
model.fit(x_train, y_train,
           batch_size=batch_size,
           epochs=epochs,
           verbose=1,
           validation_data=(x_test, y_test))
```

```
↳ Train on 4989 samples, validate on 555 samples
```

```
Epoch 1/12
4989/4989 [=====] - 703s 141ms/step - loss: 1.1764 - acc: 0.
Epoch 2/12
4989/4989 [=====] - 697s 140ms/step - loss: 0.6262 - acc: 0.
Epoch 3/12
4989/4989 [=====] - 703s 141ms/step - loss: 0.5894 - acc: 0.
Epoch 4/12
4989/4989 [=====] - 702s 141ms/step - loss: 0.5021 - acc: 0.
Epoch 5/12
4989/4989 [=====] - 704s 141ms/step - loss: 0.4299 - acc: 0.
Epoch 6/12
4989/4989 [=====] - 703s 141ms/step - loss: 0.4062 - acc: 0.
Epoch 7/12
4989/4989 [=====] - 700s 140ms/step - loss: 0.3647 - acc: 0.
Epoch 8/12
4989/4989 [=====] - 695s 139ms/step - loss: 0.3151 - acc: 0.
Epoch 9/12
4989/4989 [=====] - 696s 140ms/step - loss: 0.3187 - acc: 0.
Epoch 10/12
4989/4989 [=====] - 697s 140ms/step - loss: 0.2760 - acc: 0.
Epoch 11/12
4989/4989 [=====] - 712s 143ms/step - loss: 0.2527 - acc: 0.
Epoch 12/12
4989/4989 [=====] - 693s 139ms/step - loss: 0.2596 - acc: 0.
<keras.callbacks.History at 0x7fa26c55beb8>
```

```
#Testing the model on test set
score = model.evaluate(x_test, y_test)
print('Test loss:', score[0])
print('Test accuracy:', score[1])
```

```
↳ 555/555 [=====] - 19s 35ms/step
Test loss: 0.7425039282790176
Test accuracy: 0.790990991420574
```

